

Large-Scale and Multi-Structured Databases

Weather Prediction App

Rojan Shrestha

Alessio Simoncini

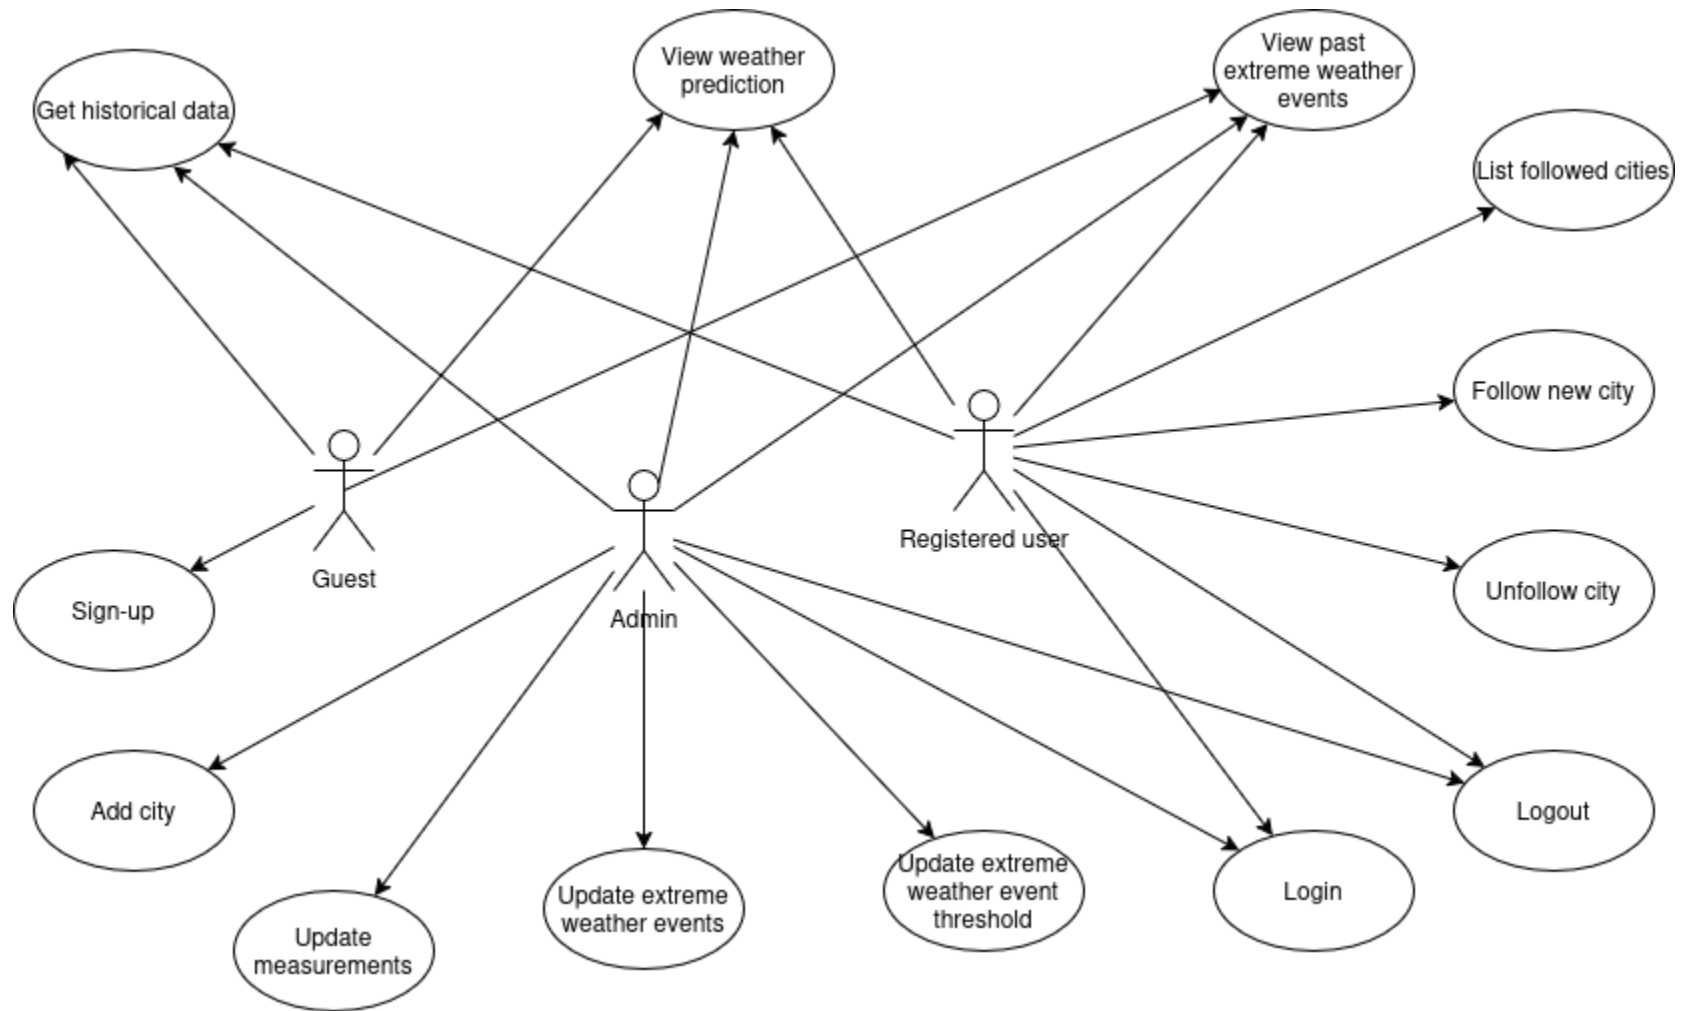
Lorenzo Vezzani

Application Highlights

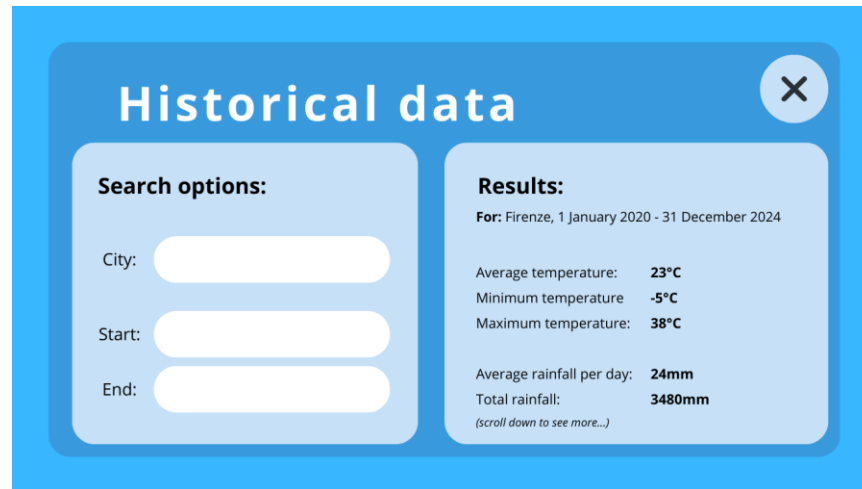
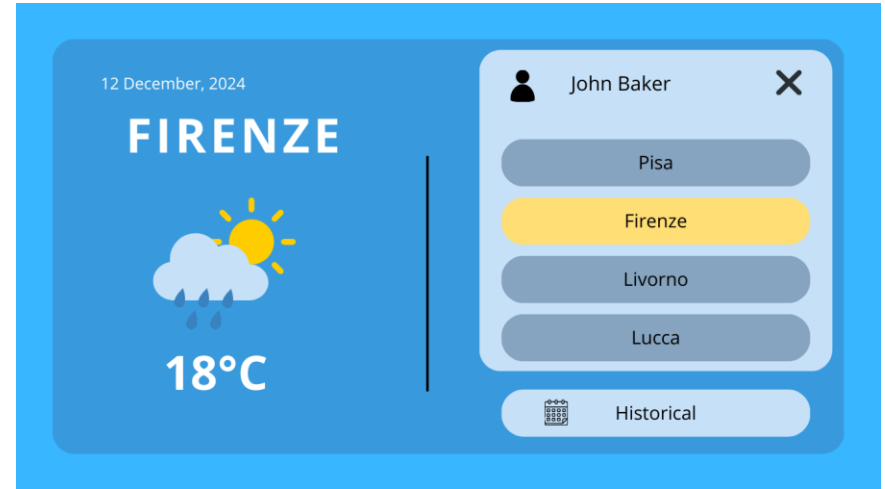
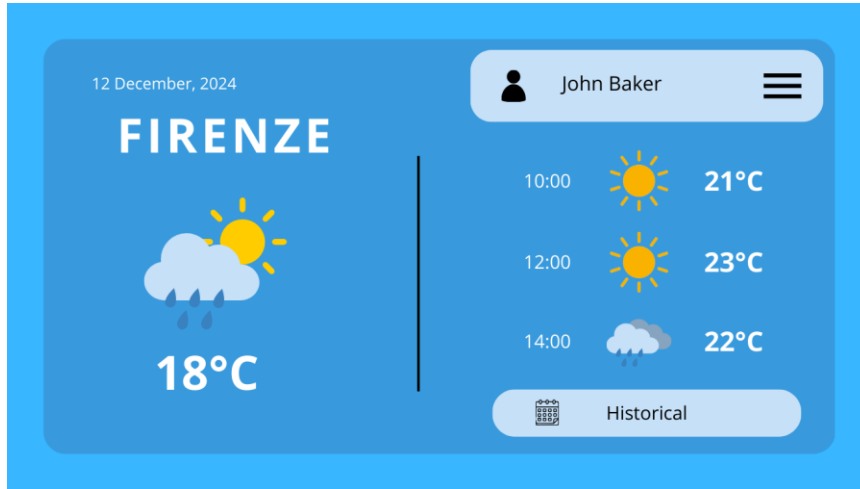
WeatherApp provides you with the weather information you need to plan your day with key features:

- **Weather Forecast** - Offers hourly, daily, or weekly weather forecasts, helping users plan ahead. Based on proximity and weather patterns the app even provides weather forecast of cities not in database.
- **Weather Trends and History** - Offers historical weather data, allowing user to analyze past weather conditions for specific cities over selected time periods.
- **Extreme Weather events** – List of extreme weather conditions that users can use to get informed about storms , hurricanes or extreme temperatures.

Actors



Main mock-ups



Dataset Description

Source:

<https://open-meteo.com/en/docs/historical-weather-api>

Open-Meteo is an open source weather API that partners with national weather services around the globe.

Description: The hourly weather data over 80 years of any major city of Italy.

Volume: 75MB

Velocity/Variability: Weather information is important only when recent. After a certain period of time it will be relevant only for statistical purposes.

Non-functional requirements

Performance: Fast Read Times and Scalability

- Minimize read time as much as possible

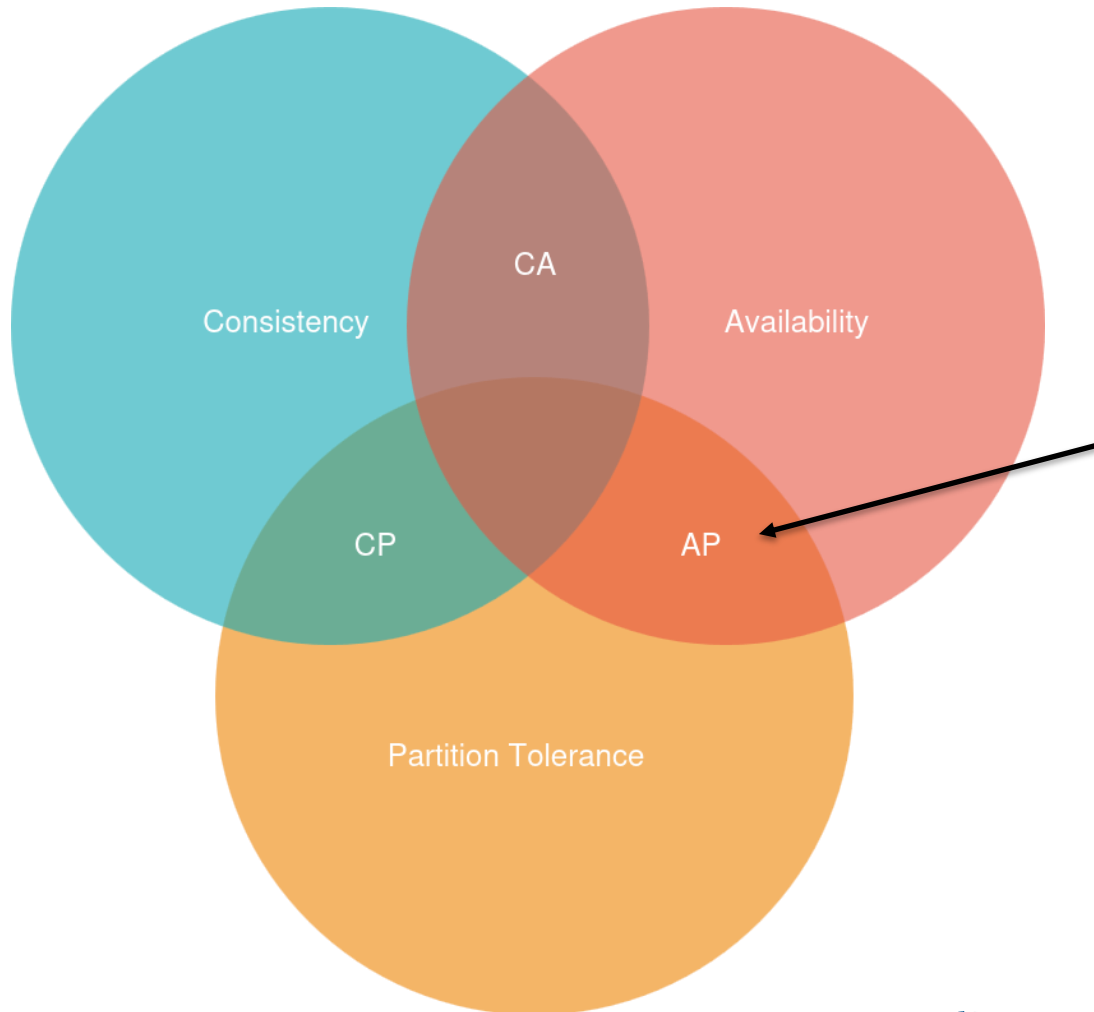
Performance: Availability

- High availability service
- Low latency even under heavy load

Security

- Authentication mechanism to ensure secure database access
- Upon login, the user receives a token. Token must be included in HTTP requests to access specific APIs

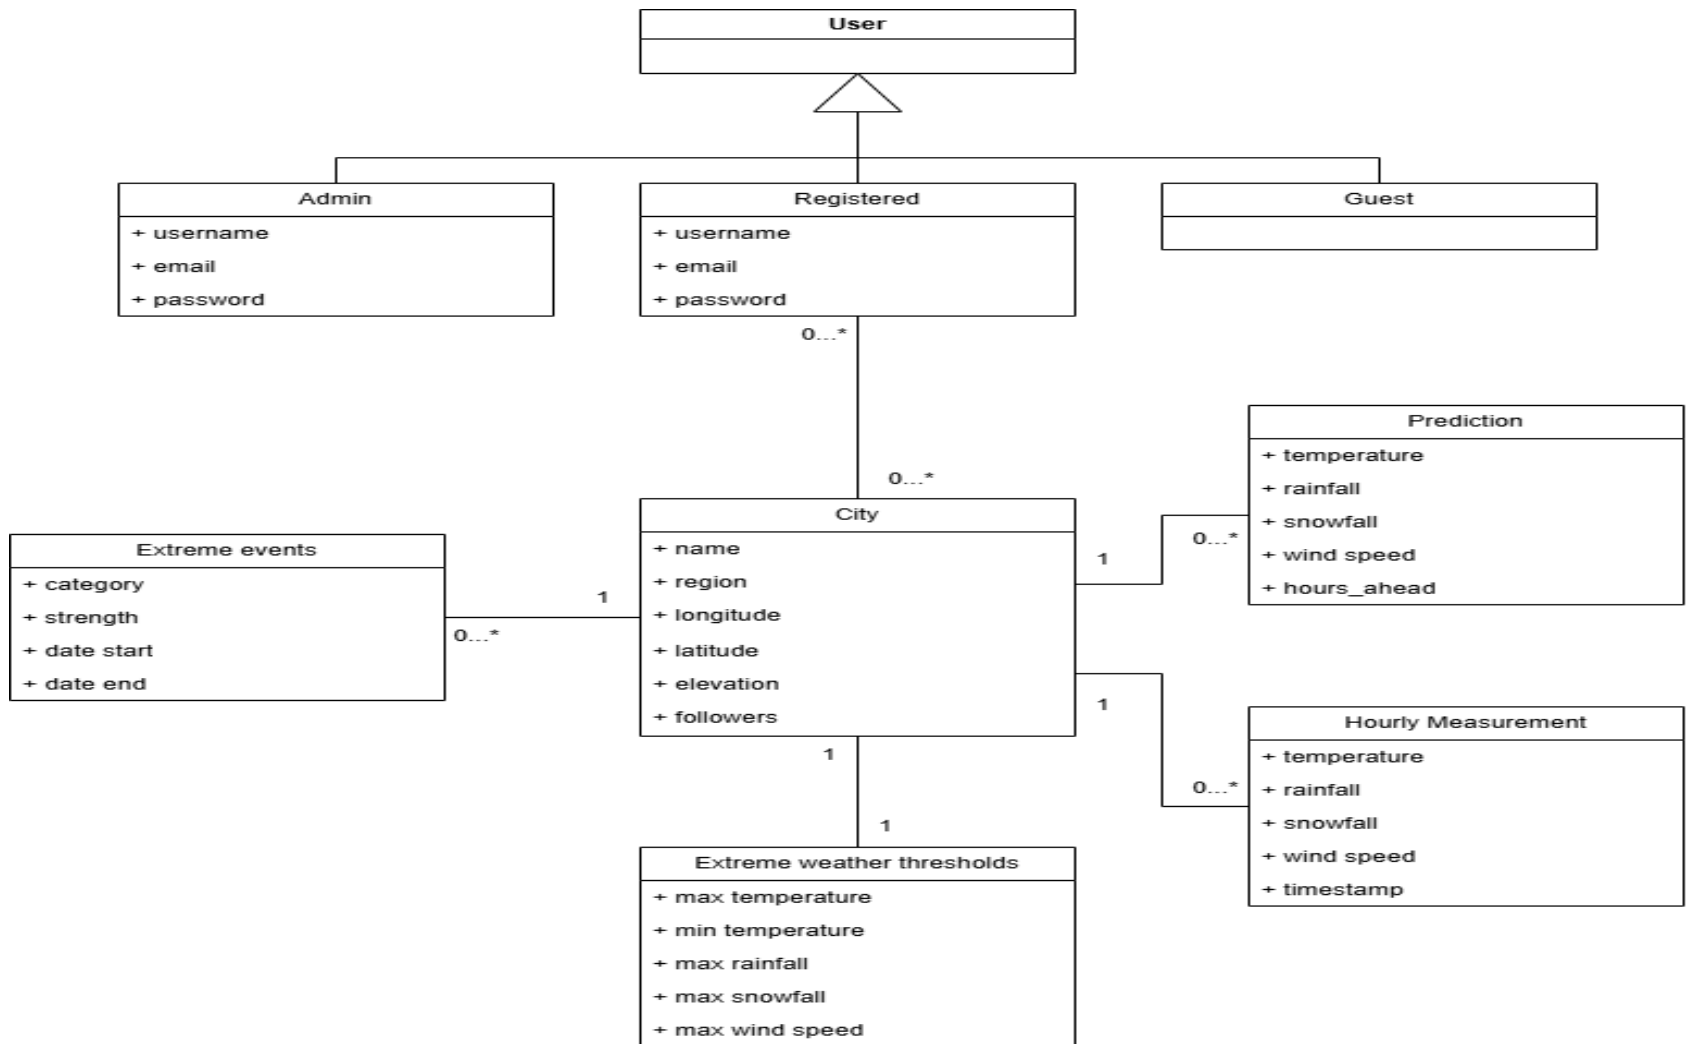
CAP theorem consideration



The application is focused on Availability and Network Partition Tolerance (AP).

- High availability to ensure fast reads, since the application is read-based
- The application is distributed, partition tolerance is not a given

UML Class Diagram



Document DB: collections and indexes

- Users

- username, password, email, listCityId

 *Data linking* strategy: stores only city IDs

Index on username for efficient login

- Cities

- name, region, coordinates, threshold, followers, city_id, eweList

 *Data embedding* used for extreme weather events (rare, read-heavy)

Default _id index (sufficient for city_id access)

- Measurements

- temperature, rainfall, snowfall, windspeed, timestamp, city_id

 *Data linking* adopted: embedding would exceed MongoDB size limits

Compound index on city_id and timestamp, expensive but essential for fast analytics over time

Experimental indexes validation

Chosen query:

```
command: {
  aggregate: 'hourly_measurements',
  pipeline: [
    {
      '$match': {
        cityId: 'tus-pis-43.7085-10.4036',
        time: {
          '$gte': ISODate('2020-05-20T00:00:00.000Z'),
          '$lte': ISODate('2020-09-20T00:00:00.000Z')
        }
      }
    },
    {
      '$group': { _id: '$cityId', avgTemperature: { '$avg': '$temperature' } }
    },
    {
      '$project': { cityId: '$_id', avgTemperature: '$avgTemperature' }
    }
  ],
  cursor: {},
  '$db': 'WeatherApp'
},
```

Experimental indexes validation

Without index:

```
executionStats: {  
  executionSuccess: true,  
  nReturned: 1,  
  executionTimeMillis: 97099,  
  totalKeysExamined: 0,  
  totalDocsExamined: 13396320,  
}
```

(The entire collection is examined)

With index:

```
{  
  v: 2,  
  key: { cityId: 1, time: 1 },  
  name: 'cityId_1_time_1',  
  unique: true  
}
```

```
executionStats: {  
  executionSuccess: true,  
  nReturned: 1,  
  executionTimeMillis: 26,  
  totalKeysExamined: 2953,  
  totalDocsExamined: 2953,  
}
```

Key-Value DB

Cities Storage

- **Key Format:**
city:{regioncode}-citycode-latitude-longitude
e.g., city:{tus}-pis-43.7085-10.4036
- **Value Type:** HASH
 - Stores useful metadata (e.g., region, elevation)

Forecast Storage

- **Key Format:**
forecast:{regioncode}-citycode-latitude-longitude:date
e.g., forecast:{tus}-pis-43.7085-10.4036:2025-06-08
- **Value Type:** STRING
 - Daily forecast including:
temperature, rainfall, snowfall, windspeed, timestamp

Redis is also used for storing user **token**: created at login and deleted at logout.

Handling intra-DB consistency

The only shared entity between **MongoDB** and **Redis** is **city**.

The critical operation is `saveCity(·)`:

- Adding the city to MongoDB.
- Storing city metadata in Redis.

Rollback Mechanism

- In case of any failure:
 - The inserted city is **deleted from MongoDB** (if added).
 - The corresponding Redis hash is **deleted** (if added).

Result

- The `saveCity(·)` operation ensures **strong intra-database consistency**.
- No partial insertions of **city** remain if the process is interrupted.

Clustering and Sharding: mongo-conf

```
mongod --replSet lsmdb --dbpath ~/data --port 27020 --bind_ip localhost,10.1.1.9 --oplogSize 200
```

```
mongod --replSet lsmdb --dbpath ~/data --port 27020 --bind_ip localhost,10.1.1.84 --oplogSize 200
```

```
mongod --replSet lsmdb --dbpath ~/data --port 27020 --bind_ip localhost,10.1.1.87 --oplogSize 200
```

```
rsconf = { _id: "lsmdb",  
  members: [  
    { _id: 0, host: "10.1.1.9:27020", priority:1},  
    { _id: 1, host: "10.1.1.84:27020", priority:2},  
    { _id: 2, host: "10.1.1.87:27020", priority:5}  
  ]};
```

Clustering and Sharding: redis-conf

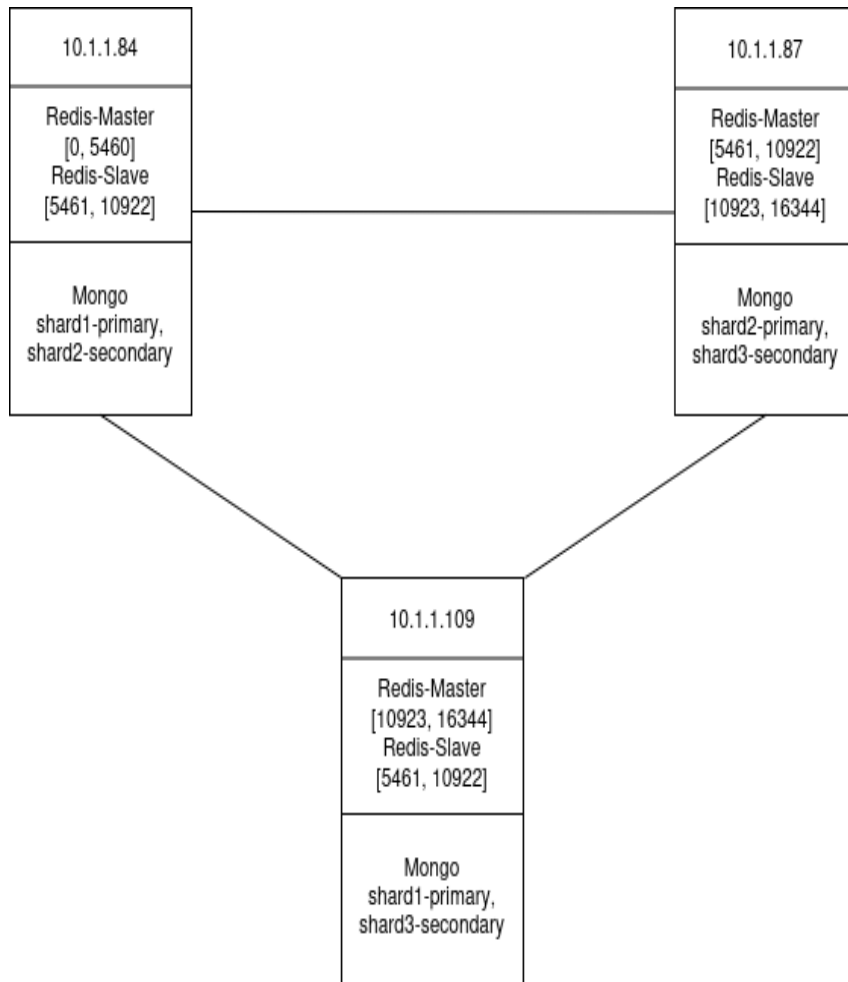
port 6379 # 6380 for the slaves
cluster-enabled yes
cluster-config-file nodes.conf
cluster-node-timeout 5000
appendonly yes
bind 0.0.0.0
maxmemory 512mb
maxmemory-policy allkeys-lfu

Runned on 10.1.1.84:6379



```
127.0.0.1:6379> cluster nodes
9c507545202c13c8ee2709514a23ac5804d44c7f 10.1.1.87:6380@16380 slave cec1a25566d82201b2e81ea06d0d3ed489c5801b 0 1750510225000 4 connected
cec1a25566d82201b2e81ea06d0d3ed489c5801b 10.1.1.84:6379@16379 myself,master - 0 1750510224000 1 connected 0-5460
086b403052c6b0ef8eb8c69259a3866bf5ec76a7 10.1.1.87:6379@16379 slave 5a676b5d82b8f0f21518f3a10992a5bd5ee0b139 0 1750510225570 7 connected
5a676b5d82b8f0f21518f3a10992a5bd5ee0b139 10.1.1.9:6380@16380 master - 0 1750510226073 7 connected 5461-10922
25b894cec319d039410c0218b4044f8596678959 10.1.1.9:6379@16379 master - 0 1750510225369 5 connected 10923-16383
582c8d338698ce297fdc54f94c8ae624e087be78 10.1.1.84:6380@16380 slave 25b894cec319d039410c0218b4044f8596678959 0 1750510226373 5 connected
```

Clustering and Sharding



VMs Available

- Deployed on 3 UniPi VMs:
 - 10.1.1.9, 10.1.1.84, 10.1.1.87
- Read-Optimized Architecture

MongoDB Clustering

- Write concern: `w=1` → availability
- Read preference: `local` → fast reads

Mongo Sharding Strategy (planned)

- `user`: shard by `_id`
- `city`: shard by `_id`
- `hourly_measurement`: shard by `cityId`

Redis Sharding

- Keys structured with region code inside `{}`:
 - `city:{tus}-pis-...`
 - `forecast:{tus}-pis-...-2025-01-01`

Swagger UI REST APIs documentation

User

- POST /user/register
- POST /user/login
- POST /user/logout

Forecasts

- GET /forecast/today
- GET /forecast/today/arbitrary-city
- GET /forecast/day
- GET /forecast/day/arbitrary-city

City

- POST /city/add
- PUT /city/update-thresholds
- GET /city/by-name
- GET /city/all
- GET /favorites
- PUT /favorites
- DELETE /favorites

Swagger UI REST APIs documentation

Extreme Weather Event (Admin)

- DELETE /ewe/duplicates/range
- DELETE /ewe/duplicates/all

Admin – Data Update

- PUT /data-manager/update/forecasts
- PUT /data-manager/update/measurements
- PUT /data-manager/update/ewes

(Some) Analytics

- GET /analytics/measurement/city/average-per-month
- GET /analytics/measurement/recent/city/total-per-day / average-per-day
- GET /analytics/ewe/strength/maximum / average
- GET /analytics/ewe/duration/longest / average
- GET /analytics/ewe/count / count-per-month / count-of-at-least-strength

Live Demo with Postman